

sheet03_Rename

November 11, 2024


From Oliver Sange, Elias Huber and Sam Rouppe van der Voort

```
[ ]: import numpy as np
      from matplotlib import pyplot as plt
      from sklearn.linear_model import LinearRegression
```

0.1 3 Visualize Regularization Contours

```
[25]: # load the data
      data = np.load('data/linreg.npz')
      x = data['X']
      y = data['Y']
      print(f'x.shape: {x.shape}, "y.shape:" {y.shape}')
      plt.scatter(*x, c=y);
      plt.colorbar()
      plt.show()
```

x.shape: (2, 100), "y.shape:" (1, 100)



```
[26]: # create a grid of points in the parameter space
b1, b2 = np.linspace(-1, 3, 101), np.linspace(-1, 3, 101)
bs = np.stack(np.meshgrid(b1, b2, indexing='ij'), axis=-1)
bs.shape
```

```
[26]: (101, 101, 2)
```

0.1.1 (a)

```
[27]: ridge_term = bs[..., 0]**2 + bs[..., 1]**2
lasso_term = np.abs(bs[..., 0]) + np.abs(bs[..., 1])

# Plotting
fig, ax = plt.subplots(1, 2, figsize=(20, 7))


# Ridge regression
contour1 = ax[0].contourf(bs[..., 0], bs[..., 1], ridge_term, cmap='viridis')
fig.colorbar(contour1, ax=ax[0])
ax[0].set_title('Contour Plot of Ridge Regularization Term')
ax[0].set_xlabel('1')
ax[0].set_ylabel('2')
```

```

# Lasso regression
contour2 = ax[1].contourf(bs[... , 0], bs[... , 1], lasso_term, cmap='viridis')
fig.colorbar(contour2, ax=ax[1])
ax[1].set_title('Contour Plot of Lasso Regularization Term')
ax[1].set_xlabel(' 1')
ax[1].set_ylabel(' 2')

plt.show()

```



0.1.2 (b)

```

[28]: # Reshape bs
b1_vals = bs[... , 0]
b2_vals = bs[... , 1]

# Compute predictions
y_pred_grid = b1_vals[:, :, np.newaxis] * x[0] + b2_vals[:, :, np.newaxis] * x[1]


# Compute residuals
residuals = y - y_pred_grid

# Compute SSQ
SSQ = np.sum(residuals**2, axis=-1)

# Plot
fig, ax = plt.subplots(figsize=(8, 6))
contour = ax.contourf(b1_vals, b2_vals, SSQ, cmap='viridis')
fig.colorbar(contour, ax=ax)
ax.set_title('SSQ as a Function of 1 and 2')
ax.set_xlabel(' 1')

```

```
ax.set_ylabel(' 2')
plt.show()
```



0.1.3 (c)

```
[29]: # TODO: for each lambda, plot both ridge regression and lasso loss functions
lambdas = [0, 10, 50, 100, 200, 300]

for lam in lambdas:
    ridge_term = lam * (bs[..., 0]**2 + bs[..., 1]**2)
    lasso_term = lam * (np.abs(bs[..., 0]) + np.abs(bs[..., 1]))

    # Plotting
    fig, ax = plt.subplots(1, 2, figsize=(20, 7))

    # Ridge regression
    contour1 = ax[0].contourf(bs[..., 0], bs[..., 1], ridge_term,
    cmap='viridis')
```

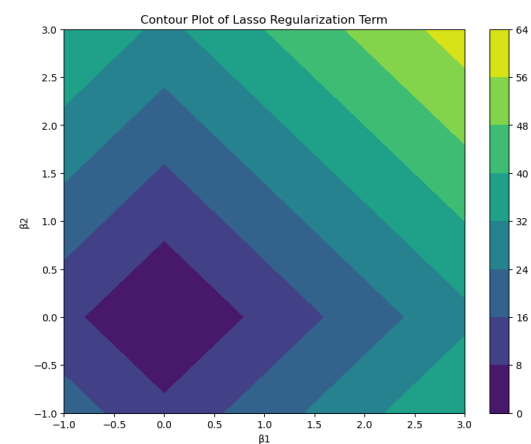
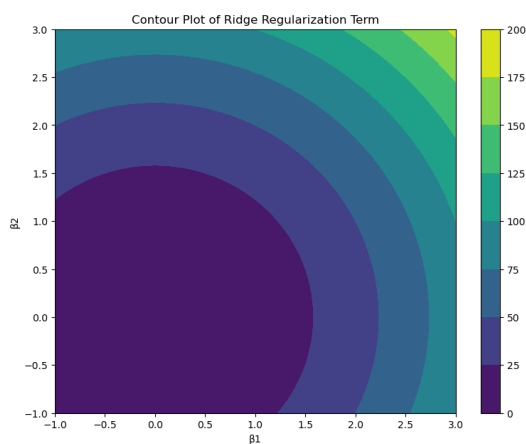
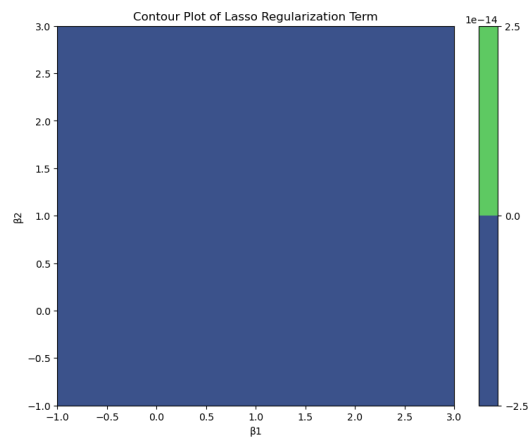
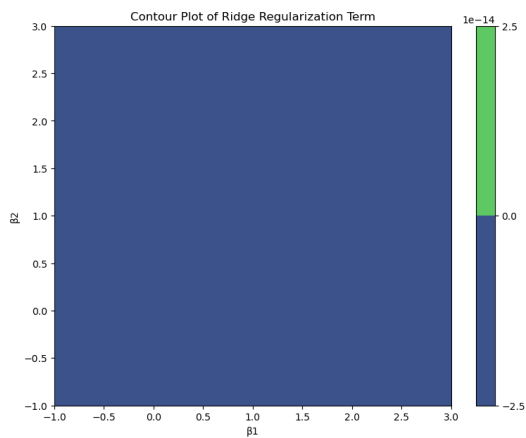
```


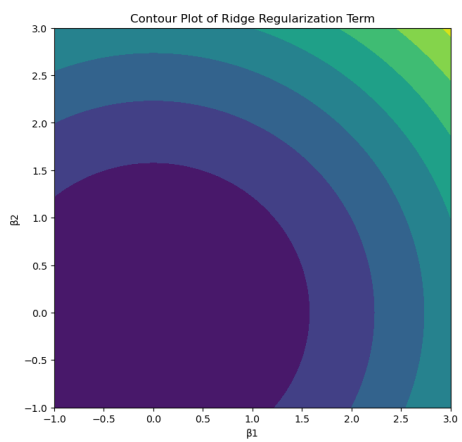

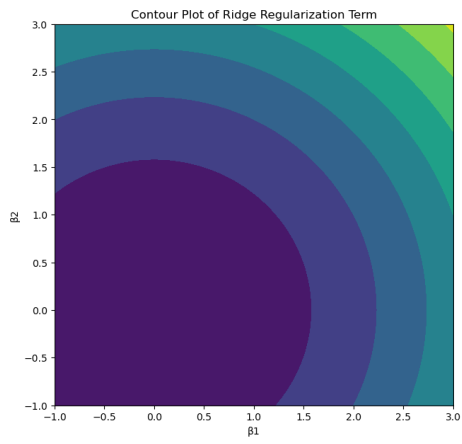


fig.colorbar(contour1, ax=ax[0])
ax[0].set_title('Contour Plot of Ridge Regularization Term')
ax[0].set_xlabel('1')
ax[0].set_ylabel('2')


# Lasso regression
contour2 = ax[1].contourf(bs[...], 0], bs[...], 1], lasso_term,
cmap='viridis')
fig.colorbar(contour2, ax=ax[1])
ax[1].set_title('Contour Plot of Lasso Regularization Term')
ax[1].set_xlabel('1')
ax[1].set_ylabel('2')

plt.show()

```







For a larger λ , the colorbar will extend to higher values. This means that the colors in the plot will represent a stronger penalty, and the “cost” of larger coefficients β_1 and β_2 becomes more pronounced in the optimization.

0.2 4 CT Reconstruction

First, set up the design matrix. (Run this once to save it to the disk)

```
[15]: # create design matrix
# don't change any of this, just run it once to create and save the design
      ↪matrix
import os

n_parallel_rays = 70
n_ray_angles = 30
res = (99, 117)
print("Number of pixels in the 2d image:", np.prod(res))
print("Total number of rays:", n_parallel_rays * n_ray_angles)

def rot_mat(angle):
    c, s = np.cos(angle), np.sin(angle)
    return np.stack([np.stack([c, s], axis=-1), np.stack([-s, c], axis=-1)],
                    ↪axis=-1)

kernel = lambda x: np.exp(-x**2/sigma**2/2)

if not os.path.exists('data/design_matrix.npy'):
    xs = np.arange(0, res[1]+1) - res[1]/2 # np.linspace(-1, 1, res[1] + 1)
    ys = np.arange(0, res[0]+1) - res[0]/2 # np.linspace(-1, 1, res[0] + 1)

    # rays are defined by origin and direction
    ray_offset_range = [-res[1]/1.5, res[1]/1.5]
```

```

n_rays = n_parallel_rays * n_ray_angles

ray_angles = np.linspace(0, np.pi, n_ray_angles, endpoint=False) + np.pi/
↪n_ray_angles

# offsets for ray_angle = 0, i.e. parallel to x-axis
ray_0_offsets = np.stack([np.zeros(n_parallel_rays), np.
↪linspace(*ray_offset_range, n_parallel_rays)], axis=-1)
ray_0_directions = np.stack([np.ones(n_parallel_rays), np.
↪zeros(n_parallel_rays)], axis=-1)

ray_rot_mats = rot_mat(ray_angles)

ray_offsets = np.einsum('oi,aij->aoj', ray_0_offsets, ray_rot_mats).
↪reshape(-1, 2)
ray_directions = np.einsum('oi,aij->aoj', ray_0_directions, ray_rot_mats).
↪reshape(-1, 2)

sigma = 1
xsc = (xs[1:] + xs[:-1]) / 2
ysc = (ys[1:] + ys[:-1]) / 2
b = np.stack(np.meshgrid(xsc, ysc), axis=-1).reshape(-1, 2)
a = ray_offsets
v = ray_directions
v = v / np.linalg.norm(v, axis=-1, keepdims=True)
p = ((b[None] - a[:, None]) * v[:, None]).sum(-1, keepdims=True) * v[:,
↪None] + a[:, None]
d = np.linalg.norm(b - p, axis=-1)
d = kernel(d)
design_matrix = d.T

np.save('data/design_matrix.npy', design_matrix)
print(f'created and saved design matrix of shape {design_matrix.shape} at
↪data/design_matrix.npy')

```

Number of pixels in the 2d image: 11583

Total number of rays: 2100

```

[16]: sino = np.load('data/sino.npy')

print(f'sino shape: {sino.shape}')


# visualize sinogram as image
n_parallel_rays = 70
n_angles = 30

```



```
plt.imshow(sino.reshape(n_angles, n_parallel_rays), origin='lower')
# plt.colorbar()
plt.xlabel('parallel rays')
plt.ylabel('angles from 0 to 180 deg')
plt.show();
```

sino shape: (1, 2100)



0.2.1 (a)

```
[17]: design_matrix = np.load('data/design_matrix.npy')
      numbers = np.random.randint(0, design_matrix.shape[1], size = 4)
      print(f"shape design matrix: {design_matrix.shape}")


      image1 = design_matrix[:, numbers[0]].reshape(res)
      image2 = design_matrix[:, numbers[1]].reshape(res)
      image3 = design_matrix[:, numbers[2]].reshape(res)
      image4 = design_matrix[:, numbers[3]].reshape(res)



      plt.imshow(image1, origin='lower')
      plt.show()
      plt.imshow(image2, origin='lower')
      plt.show()
      plt.imshow(image3, origin='lower')
      plt.show()
      plt.imshow(image4, origin='lower')
      plt.show()
```


```
# Each image shows the trajectory and its corresponding intensity through the ↵  
↵ final image
```

```
# TODO: visualize four random columns as images, using an image shape of (99, ↵  
↵ 117)
```

shape design matrix: (11583, 2100)









0.2.2 (b)



```
[18]: # TODO: visualize four random rows as images, using an images
numbers = np.random.randint(0,design_matrix.shape[0],size = 4)
res2 = (30,70)

image1 = design_matrix[numbers[0]].reshape(res2)
image2 = design_matrix[numbers[1]].reshape(res2)
image3 = design_matrix[numbers[2]].reshape(res2)
image4 = design_matrix[numbers[3]].reshape(res2)

plt.imshow(image1, origin='lower')
plt.show()
plt.imshow(image2, origin='lower')
plt.show()
plt.imshow(image3, origin='lower')
plt.show()
plt.imshow(image4, origin='lower')
```

```
plt.show()  
# The images map the readout intensity of a detector to the corresponding ray  
↔ and angle
```






0.2.3 (c)

```
[19]: # TODO: solve the reconstruction with linear regression and visualize the result
print(f"shapes: sino {sino.shape} design_matrix: {design_matrix.shape}")
# project components onto sino
Y = sino.squeeze()
reg = LinearRegression().fit(design_matrix.T, Y)
print(reg.score(design_matrix.T, Y))
prediction = reg.predict(design_matrix.T)
new_sino = prediction.reshape((30, 70))
plt.imshow(new_sino)
plt.show()
# double check if sinogram fits
```


```
image = reg.coef_.reshape(res)
plt.imshow(image)
```

we observe that the reconstructed image is not interpretable

```
shapes: sino (1, 2100) design_matrix: (11583, 2100)
0.99999999900717896
```



```
[19]: <matplotlib.image.AxesImage at 0x7f5ce8aa6e80>
```



```
[11]: # TODO: solve the reconstruction with ridge regression and visualize the result
# Optional: try out different regularization strengths and observe the influence
from sklearn.linear_model import Ridge

reg = Ridge(alpha=0.1).fit(design_matrix.T, Y)
print(reg.score(design_matrix.T, Y))
prediction = reg.predict(design_matrix.T)
image = reg.coef_.reshape(res)
plt.imshow(image)
plt.show()

reg = Ridge(alpha=1).fit(design_matrix.T, Y)
print(reg.score(design_matrix.T, Y))
prediction = reg.predict(design_matrix.T)
image = reg.coef_.reshape(res)
plt.imshow(image)
plt.show()

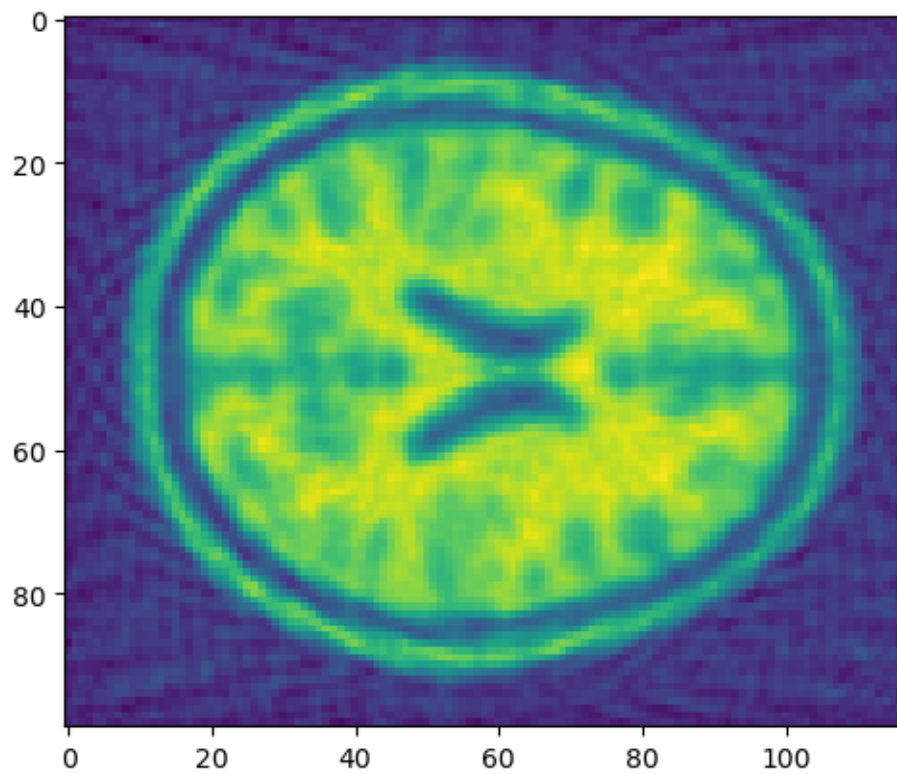
reg = Ridge(alpha=10).fit(design_matrix.T, Y)
print(reg.score(design_matrix.T, Y))
prediction = reg.predict(design_matrix.T)
image = reg.coef_.reshape(res)
```




```
plt.imshow(image)
plt.show()
```

We observe that $\alpha = 0.1$ delivers the best image quality as well as the highest prediction score


0.9999999710587341



0.9999997054099179





0.9999783629561321



0.3 5 Bonus: X-Ray Free-Electron Lasers

```
[9]: sino = np.load('data/sino.npy').reshape(n_angles, n_parallel_rays)
plt.imshow(sino)
plt.title('original sinogram')
plt.show()

order = np.arange(n_angles)
np.random.shuffle(order)
sino_shuffled = sino[order]
plt.imshow(sino_shuffled)
plt.title('shuffled sinogram')
plt.show()
```



[]:

Machine Learning and Physics Sheet 05

a) $\hat{\beta}_{\text{ridge}} = \underset{\beta}{\operatorname{argmin}} \|y - X^T \beta\|_2^2 + \lambda \|\beta\|_2^2$

$$= \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^n \left(y_i - (\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2}) \right)^2 + \lambda (\beta_0^2 + \beta_1^2 + \beta_2^2)$$

We then get $\hat{\beta}_{\text{ridge}} = (XX^T + \lambda I)^{-1} Xy$

→ Since the regularization penalizes any values of β_i ,
This results in a bias for the intercept β_0
towards values close to 0.

b) We modify: $\hat{\beta}_{\text{ridge, mod}} = \underset{\beta}{\operatorname{argmin}} \|y - X^T \beta\|_2^2 + \lambda \|\beta\|_2^2 - \lambda \beta_0^2$
to remove the punishment of the β_0 component.

c) Since we remove the dependency on β_0 of the penalty term,
The total penalty losses read:

$\lambda(\beta_1^2 + \beta_2^2)$ for b) and $\lambda(\beta_0^2 + \beta_1^2 + \beta_2^2)$ for a).

If we now assume we find a minimal value, the regularization contour (sets of parameters with same value) then

solves the equation

$\lambda(\beta_1^2 + \beta_2^2) = L$ for b) and $\lambda(\beta_0^2 + \beta_1^2 + \beta_2^2) = L$ for a).

We thus have the surface of a sphere for a)
and the surface of a cylinder for b).

(Since for all values of β_0 we have the same penalty).

$$\textcircled{2} a) y_n = \beta^T x_n + \varepsilon_n \quad \text{with} \quad \varepsilon_n \sim \mathcal{N}(0, \sigma^2)$$

$$\begin{aligned} \hat{\beta} &= \underset{\beta}{\operatorname{argmax}} \sum_{n=1}^N \log \mathcal{N}(y_n | \beta^T x_n, \sigma^2) \\ &= \underset{\beta}{\operatorname{argmax}} \sum_{n=1}^N \log e^{-(y_n - x_n^T \beta)^T (y_n - x_n^T \beta) / \sigma^2} \\ &= \underset{\beta}{\operatorname{argmax}} \log \prod_{n=1}^N e^{-(y_n - x_n^T \beta)^T (y_n - x_n^T \beta) / \sigma^2} \\ &= \underset{\beta}{\operatorname{argmax}} - \sum_{n=1}^N (y_n - x_n^T \beta)^T (y_n - x_n^T \beta) / \sigma^2 \\ &= \underset{\beta}{\operatorname{argmin}} \sum_{n=1}^N (y_n - x_n^T \beta)^T (y_n - x_n^T \beta) \end{aligned}$$

$$b) \frac{1}{\sigma^2} = \underset{\sigma^2}{\operatorname{argmax}} \sum_{n=1}^N \log \mathcal{N}(y_n | \hat{\beta}^T x_n, \sigma^2)$$

$$\begin{aligned} a) &= \underset{\sigma^2}{\operatorname{argmin}} \sum_{n=1}^N (y_n - x_n^T \hat{\beta})^T (y_n - x_n^T \hat{\beta}) / \sigma^2 & \hat{\beta} &= (X^T X)^{-1} X^T y \\ &= \underset{\sigma^2}{\operatorname{argmin}} (y - (X^T (X^T X)^{-1} X) y)^T (y - X^T (X^T X)^{-1} X y) / \sigma^2 & y &= \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix} \end{aligned}$$